



SINCE TOMORROW

Event-Driven cloud Architecture, CQRS and Event Sourcing for User Management

2019/08/09

About

Company	Molo17 S.r.l.
Author	Lorenzo Busin
State	Approved
Use	Internal
Email	lorenzo.busin@gmail.com

Description

Documentation about cloud architectures that use an Event-Driven approach and implement CQRS and Event Sourcing for User Management.

Table of contents

1	Introduction	4
1.1	Project purpose	4
1.2	References	4
2	Amazon Web Services	5
2.1	Lambda	5
2.1.1	Write mode Lambda functions	5
2.1.1.1	PushOperationAggregateToSQS	5
2.1.1.2	CommandOperationAggregate	6
2.1.1.3	Mediator	8
2.1.1.4	OperationAggregate	9
2.1.1.5	Recovery	10
2.1.2	Read mode Lambda functions	11
2.1.2.1	ReadOperationAggregate	11
2.1.3	CloudWatch Logs	12
2.2	DynamoDB	13
2.3	API Gateway	13
2.4	Authorizer	13
2.5	Simple Queue Service	14
3	User Management	15
3.1	Aggregates	15
3.1.1	User	15
3.1.2	Role	15
3.1.3	Authorization	15
3.1.4	Group	15
3.2	Admin side	16
3.2.1	Authentication	16
3.2.2	Use cases	16
3.3	User side	17
3.3.1	Authentication	17
3.3.2	Use cases	18
4	Serverless Framework	20
4.1	Description	20
4.2	Serverless.yml	20
4.2.1	Serverless configuration for user management	20
4.3	Project's root	23
5	CQRS	24
5.1	Architecture overview	25
5.2	Write model	26
5.3	Read model	26
6	Event Sourcing	27

6.1	Event store	27
6.1.1	Event object	27
7	Extension points	29
7.1	New aggregates	29
7.2	New write operation	29
7.3	New read operation	29
8	Setup	31
8.1	Installing Node.js	31
8.2	Installing Serverless Framework	31
8.3	Setup AWS credentials	31
8.4	Deploying the serverless User Management service	32

List of figures

1	Fetch POST API	5
2	Lambda: pushOperationAggregateToSQS	6
3	Lambda: commandOperationAggregate	7
4	DynamoDB event object	8
5	Parsed DynamoDB event object	8
6	Lambda: mediator	9
7	Lambda: operationAggregate	10
8	Lambda: recovery	11
9	Fetch GET API	11
10	Lambda: readOperationAggregate	12
11	Authorizer workflow	14
12	Admin use cases	16
13	User use cases	18
14	Service, plugins and custom - serverless.yml	20
15	Provider - serverless.yml	21
16	Custom gateway response - serverless.yml	21
17	DynamoDB table - serverless.yml	22
18	Lambda triggered by SQS event - serverless.yml	22
19	Lambda triggered by GET endpoint - serverless.yml	23
20	Project's root	23
21	Architecture overview	25
22	Write model sequence diagram	26
23	Read model sequence diagram	26
24	Event object item	28
25	Node.js download page	31
26	Setup user details - AWS credentials	32
27	Setup policy - AWS credentials	32

1 Introduction

1.1 Project purpose

This project shows how to build an Event-Driven cloud architecture that apply the CQRS pattern and Event Sourcing for user management, showing how it was implemented and how to extend it with new feature or new aggregates.

There are two side of the application:

- Admin side: an administrator can execute CRUD operation for each aggregate(users, roles, authorizations and groups) and rebuild the system from a specific timestamp;
- User side: a user can sign in into the application the first time logs in; then the user can see or update its profile information.

The application also integrates the authentication function managed by a third-party provider named Auth0.

1.2 References

- CQRS - Martin Fowler:
martinfowler.com/bliki/CQRS.html;
- Event Sourcing - Martin Fowler:
martinfowler.com/eaaDev/EventSourcing.html;
- AWS Lambda docs:
<https://docs.aws.amazon.com/lambda/index.html>;
- AWS DynamoDB docs:
<https://docs.aws.amazon.com/dynamodb/index.html>;
- AWS API Gateway docs:
<https://docs.aws.amazon.com/apigateway/index.html>;
- AWS SQS docs:
<https://docs.aws.amazon.com/sqs/index.html>;
- Serverless Framework:
<https://serverless.com/>;
- Auth0 docs:
<https://auth0.com/docs>.

2 Amazon Web Services

Amazon Web Services (AWS) is a subsidiary of Amazon that provides on-demand cloud computing platforms to individuals, companies, and governments, on a pay-as-you-go basis. These cloud computing web services provide a set of primitive abstract technical infrastructure and distributed computing building blocks and tools. AWS version of virtual computers emulate most of the attributes of a real computer, including hardware central processing units, graphics processing units, memories and storage.

2.1 Lambda

AWS Lambda is an Event-Driven serverless computing platform provided by Amazon. It is a computing service that runs code in response to events and automatically manages the computing resources required by that code. The purpose of Lambda is to simplify building applications that are responsive to events and new information. AWS targets starting a Lambda instance within milliseconds of an event.

In this project, lambda functions are the computing core for the execution of all commands and are written in Node.js.

2.1.1 Write mode Lambda functions

2.1.1.1 PushOperationAggregateToSQS

Those functions starts the flow and are triggered when fetching the corresponding API gateway URL with a POST request like this:

```
fetch(linkCreateUserAPI_POST, {
  method: 'post',
  headers: {
    'Content-Type': 'application/json',
    'Authorization': 'Bearer ' + localStorage.getItem('id_token')
  },
  body: JSON.stringify({
    "attribute1": "value1",
    "attribute2": "value2"
  })
}).catch(function(error){
  showError(error);
});
```

Fig. 1: Fetch POST API

Once triggered, these functions retrieve the event and put it in the *MessageBody* parameter. The event object doesn't change, it just stringified before sending the message to corresponding SQS queue.

Example:

```
module.exports.pushCreateUserToSQS = async (event, context, callback) => {
  const utils = require('./utils.js');
  const params = {
    MessageBody: JSON.stringify(event),
    QueueUrl:
      "https://sqs.eu-central-1.amazonaws.com/582373673306/createUserQueue"
  };
  const res = await utils.pushToSQS(params);
  callback(null, res);
};
```

Fig. 2: Lambda: pushOperationAggregateToSQS

2.1.1.2 CommandOperationAggregate

These functions validate the values of the attributes before storing the event in the *eventStore* and are triggered when a new event arrives to the corresponding queue.

If you have to check for duplicated attributes in the database, the function must be marked *async* because it has to wait for the result of the check operation.

In this example the function checks for a duplicated *userId* or email address and if the attributes value is empty, otherwise the event is stored.

```
module.exports.commandCreateUser = async (event, context, callback) => {
  const utils = require('./utils.js');
  const stringedEvent = event.Records[0].body.toString('utf-8');
  const eventParsed = JSON.parse(stringedEvent);
  const stringedBody = JSON.stringify(eventParsed.body);
  const eventToCheck = JSON.parse(stringedBody);
  const checkIdParams = { //params to check for duplicated userId
    TableName: 'user',
    ProjectionExpression: "userId",
    FilterExpression: "userId = :checkId",
    ExpressionAttributeValues: {
      ":checkId": eventToCheck.userId
    }
  };
  const userIdAlreadyExists = await utils.asyncCheckScanDB(checkIdParams);
  if(userIdAlreadyExists)
    callback(null, "userId already exists");
  const checkEmailParams = { //params to check for duplicated email
    TableName: 'user',
    ProjectionExpression: "email",
    FilterExpression: "email = :checkEmail",
  };
```

```
    ExpressionAttributeValues: {
      ":checkEmail": eventToCheck.email
    }
  };
  const emailAlreadyExists = await utils.asyncCheckScanDB(checkEmailParams);
  if(emailAlreadyExists)
    callback(null, "Email already exists");
  if(eventToCheck.userId == "" || eventToCheck.firstName == "" ||
    eventToCheck.lastName == "" || eventToCheck.date == "" ||
    eventToCheck.role == "" || eventToCheck.group == ""){
    callback(null, "Empty attributes");
  }
  else{
    utils.storeEvent("user", "executeCreateUserQueue", eventToCheck);
    callback(null, "User event stored");
  }
};
```

Fig. 3: Lambda: commandOperationAggregate

2.1.1.3 Mediator

This function catch the *DynamoDB* event and is triggered when a change occurs in a certain table.

You have to be careful that the *DynamoDB* events are mapped using the char type attribute value. This is an example:

```
{
  "eventId": {
    "S": "bf6dffb9-72d8-ae5f-21fa-56dd6a26d572"
  },
  "payload": {
    "M": {
      "auth": {
        "S": "{\n\t\"Authorizations\": [\n\t\t\"FullAccess\" ] }"
      },
      "roleId": {
        "S": "63c471c1-e5c7-09d0-ea8a-f20b27a0575c"
      },
      "name": {
        "S": "Admin"
      },
      "desc": {
        "S": "Full access to all resources"
      }
    }
  },
  "aggregate": {
    "S": "role"
  },
  "executionQueue": {
    "S": "executeCreateRoleQueue"
  },
  "timestamp": {
    "N": "1563358792450"
  }
}
```

Fig. 4: DynamoDB event object

You can use the *"AWS.DynamoDB.Converter"* module to parse a DynamoDB event object.

```
{
  "eventId": "bf6dffb9-72d8-ae5f-21fa-56dd6a26d572",
  "payload": {
    "auth": "{\n\t\"Authorizations\": [\n\t\t\"FullAccess\" ] }",
    "roleId": "63c471c1-e5c7-09d0-ea8a-f20b27a0575c",
    "name": "Admin",
    "desc": "Full access to all resources"
  },
  "aggregate": "role",
  "executionQueue": "executeCreateRoleQueue",
  "timestamp": 1563358792450
}
```

Fig. 5: Parsed DynamoDB event object

After that, the mediator retrieves the *executionQueue* parameter from the event object, the payload event is passed with the *MessageBody* and then sends the execution message to corresponding SQS queue.

```
module.exports.mediator = (event, context, callback) => {
  const AWS = require('aws-sdk');
  const SQS = new AWS.SQS();
  const parser = AWS.DynamoDB.Converter; //module to parse dynamodb objects
  try{
    var parsedEvent = parser.unmarshall(event.Records[0].dynamodb.NewImage);
  }catch (err) {
    console.log(err);
    callback(null, err);
  }
  const params = { //get the SQS params
    MessageBody: JSON.stringify(parsedEvent.payload),
    QueueUrl: "https://sqs.eu-central-1.amazonaws.com/582373673306/" +
      parsedEvent.executionQueue
  };
  SQS.sendMessage(params, function(err,data){ //push to SQS
    if(err){
      console.log(err);
      callback(null, err);
    }
    else
      callback(null, "Execution event pushed to SQS");
  });
};
```

Fig. 6: Lambda: mediator

2.1.1.4 OperationAggregate

These functions execute a single operation using the event payload and are triggered when a new event arrives to corresponding execution queue.

Example:

```
module.exports.createUser = async (event, context, callback) => {
  const AWS = require('aws-sdk');
  const dynamoDb = new AWS.DynamoDB.DocumentClient();
  const stringedBody = event.Records[0].body.toString('utf-8');
  const parsedBody = JSON.parse(stringedBody);
  const params = {
    TableName: 'user',
    Item: parsedBody
  };
  await dynamoDb.put(params, (err, data) => {
    if (err){
```

```

        console.log(err);
        callback(null, err);
    }
    else
        callback(null, "User created");
    }).promise();
};

```

Fig. 7: Lambda: operationAggregate

2.1.1.5 Recovery

This function allows you to rebuild the state of the system starting from a given timestamp by replaying all the events which were stored after that time into the *eventStore* table. Is important to re-execute the events one by one and in the correct order (from the oldest one); for this purpose the recovery function implements a sorting algorithm that retrieves an array of events and sorts them by timestamp, and an *async* function which sends every event to the corresponding execution queue.

```

module.exports.recovery = (event, context, callback) => {
    const AWS = require('aws-sdk');
    const dynamoDb = new AWS.DynamoDB.DocumentClient();
    const utils = require('./utils.js');
    const queryParams = {
        TableName: 'eventStore',
        ExpressionAttributeNames: {
            "#eventtimestamp": "timestamp", //timestamp is a reserved keyword
            "#eventaggregate": "aggregate" //aggregate is a reserved keyword
        },
        ProjectionExpression: "#eventtimestamp, #eventaggregate, payload, executionQueue",
        FilterExpression: "#eventtimestamp >= :timest",
        ExpressionAttributeValues: {
            ":timest": parseInt(event.body.timestamp, 10)
        }
    };
};
dynamoDb.scan(queryParams, (err, data) => {
    if (err)
        callback(null, err);
    else {
        if (data.Count == 0)
            callback(null, "Events not found");
        else {
            const stringedData = JSON.stringify(data);
            const parsedData = JSON.parse(stringedData);
            var events = parsedData.Items;

```

```
        events.sort((a, b) => { //order events by timestamp from the oldest
        var a1 = a.timestamp, b1 = b.timestamp;
        if (a1 < b1) return -1;
        if (a1 > b1) return 1;
        return 0;
        });
        utils.asyncPushToExecutionQueue(events);
        callback(null, "Recovering...");
    }
}
});
};
```

Fig. 8: Lambda: recovery

2.1.2 Read mode Lambda functions

2.1.2.1 ReadOperationAggregate

These functions work in the read side of the architecture; they query the database to retrieve the information needed without passing through a SQS queue or a mediator. This kind of events aren't stored into the *eventStore* because they don't change the current system state; these Lambda functions fetch the URL of a GET endpoint using a GET request and listen for a response result.

```
fetch(linkUserAPI_GET, {
  method: "get",
  headers: {
    'Content-Type': 'application/json',
    'Authorization': 'Bearer ' + localStorage.getItem('id_token')
  }
}).then(function(response){
  return response;
}).catch(function(error){
  showError(error);
});
```

Fig. 9: Fetch GET API

Example:

```
module.exports.readUser = (event, context, callback) => {
  const AWS = require('aws-sdk');
  const dynamoDb = new AWS.DynamoDB.DocumentClient();
  const stringedEvent = JSON.stringify(event);
  const parsedEvent = JSON.parse(stringedEvent);
  const params = { //get user by userId
    TableName: 'user',
    Key: {
      "userId": parsedEvent.userId
    },
    KeyConditionExpression: "userId = :id",
    ExpressionAttributeValues: {
      ":id": parsedEvent.userId
    }
  };
  dynamoDb.get(params, (err, data) => {
    const stringedData = JSON.stringify(data);
    if (err)
      callback(null, err);
    else{
      if(data.Count == 0)
        callback(null, "User not found");
      else {
        const response = {
          statusCode: 200,
          headers: {
            'Content-Type': 'application/json',
            'Access-Control-Allow-Origin': '*',
            'Access-Control-Allow-Credentials': true
          },
          body: stringedData
        };
        callback(null, response);
      }
    }
  });
};
```

Fig. 10: Lambda: readOperationAggregate

2.1.3 CloudWatch Logs

Amazon CloudWatch is a monitoring and management service built for developers, system operators, site reliability engineers and IT managers. CloudWatch provides you with data and actionable insights to monitor your applications, understand and respond to system-wide performance changes, optimize resource utilization, and get a unified view of operational

health. CloudWatch collects monitoring and operational data in the form of logs, metrics, and events.

This tool can be very helpful to debug your Lambda functions and to understand what happens to your system.

2.2 DynamoDB

Amazon DynamoDB is a fully managed proprietary NoSQL database service that supports key-value and document data's structures and is offered by Amazon Web Services. In this project the DynamoDB instance is used for two purpose: it is used to store events and to keep updated the aggregates views.

The tables are the following:

- **eventStore**: to keep track of the occurred events;
- **user**: to store users info;
- **role**: to store roles info;
- **authorization**: to store authorizations info;
- **group**: to store groups info.

2.3 API Gateway

Amazon API Gateway is a fully managed service that makes it easy for developers to create, publish, maintain, monitor, and secure API at any scale. You can create REST and WebSocket API that act as a “front door” for applications to access data, business logic, or functionality from your backend services.

In this project, API endpoints starting the execution flow of a write mode function through a POST request, pushing a new event into the corresponding queue; in fact you have to provide an API endpoint for each function. On the other side, to query the database you have to reach the endpoint through a GET request to get a result response.

2.4 Authorizer

A Lambda authorizer (also known as a custom authorizer) is an API Gateway feature that uses a Lambda function to control access to your API.

A Lambda authorizer is useful if you want to implement a custom authorization scheme that uses a bearer token authentication strategy or that uses request parameters to determine the caller's identity.

When a client makes a request to one of your API's methods, API Gateway calls your Lambda authorizer, which takes the caller's identity as input and returns an IAM policy as output.

There are two types of Lambda authorizers:

- A token-based Lambda authorizer (also called TOKEN authorizer) receives the caller's identity in a bearer token, such as a JSON Web Token (JWT);

- A request parameter-based Lambda authorizer (also called a REQUEST authorizer) receives the caller's identity in a combination of headers, query string parameters, stage variables and context variables.

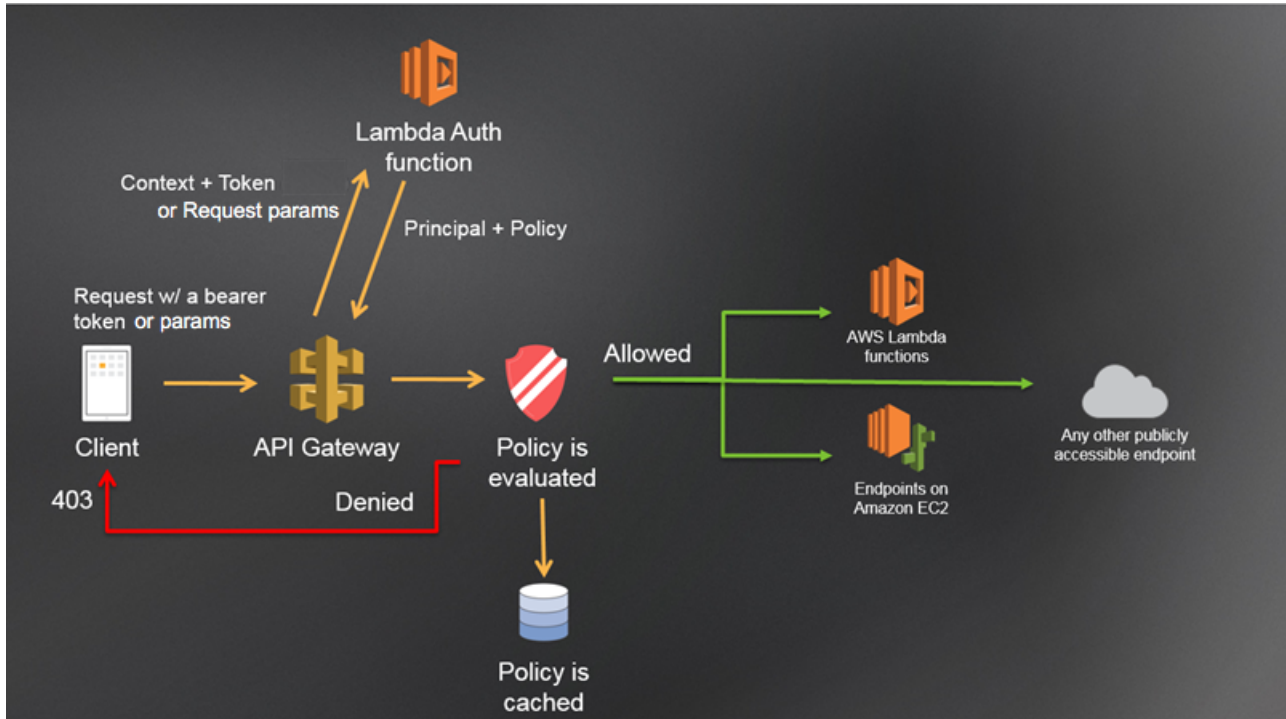


Fig. 11: Authorizer workflow

In this project the authorizer is used to protect the access to public endpoints. There are two different authorizer functions, one for admins and one for users. The workflow is the same, the only difference refers to the *client id* and the *public key* provided by *Auht0*, which are needed to authenticate the client identity.

2.5 Simple Queue Service

Simple Queue Service(SQS) is a distributed message queuing service introduced by Amazon. It supports programmatic sending of messages via web service applications as a way to communicate over the Internet. SQS is intended to provide a highly scalable hosted message queue that resolves issues arising from the common producer-consumer problem or connectivity between producer and consumer.

In this project, for each function you have to use two queues:

- **OperationAggregateQueue:** this queues receive new events from an API endpoint and trigger the corresponding *commandOperationAggregate* function;
- **ExecuteOperationAggregateQueue:** this queues receive new events from the *mediator* function and trigger the corresponding *operationAggregate* function.

3 User Management

3.1 Aggregates

3.1.1 User

- **UserId**: this attribute must be unique and refers to the *user_id* key in the corresponding *Auth0* users table(without first 6 characters 'Auth0|');
- **FirstName**: user's first name;
- **LastName**: user's last name;
- **Date**: user's birth date;
- **Email**: this attribute must be unique and respect the right format;
- External links:
 - **Role**: user's role;
 - **Group**: user's belonging group.

3.1.2 Role

- **RoleId**: is a UUID;
- **Name**: this attribute must be unique;
- **Desc**: role's description;
- External links:
 - **Auth**: JSON array of role's authorizations.

3.1.3 Authorization

- **AuthId**: is a UUID;
- **Name**: this attribute must be unique;
- **Desc**: auth's description.

3.1.4 Group

- **GroupId**: is a UUID;
- **Name**: this attribute must be unique;
- **Desc**: group's description.

3.2 Admin side

3.2.1 Authentication

The authentication feature is provided by a third-party provider, Auth0. In the admin side only administrators can log in; the client must be signed into the corresponding application's database on Auth0.

3.2.2 Use cases

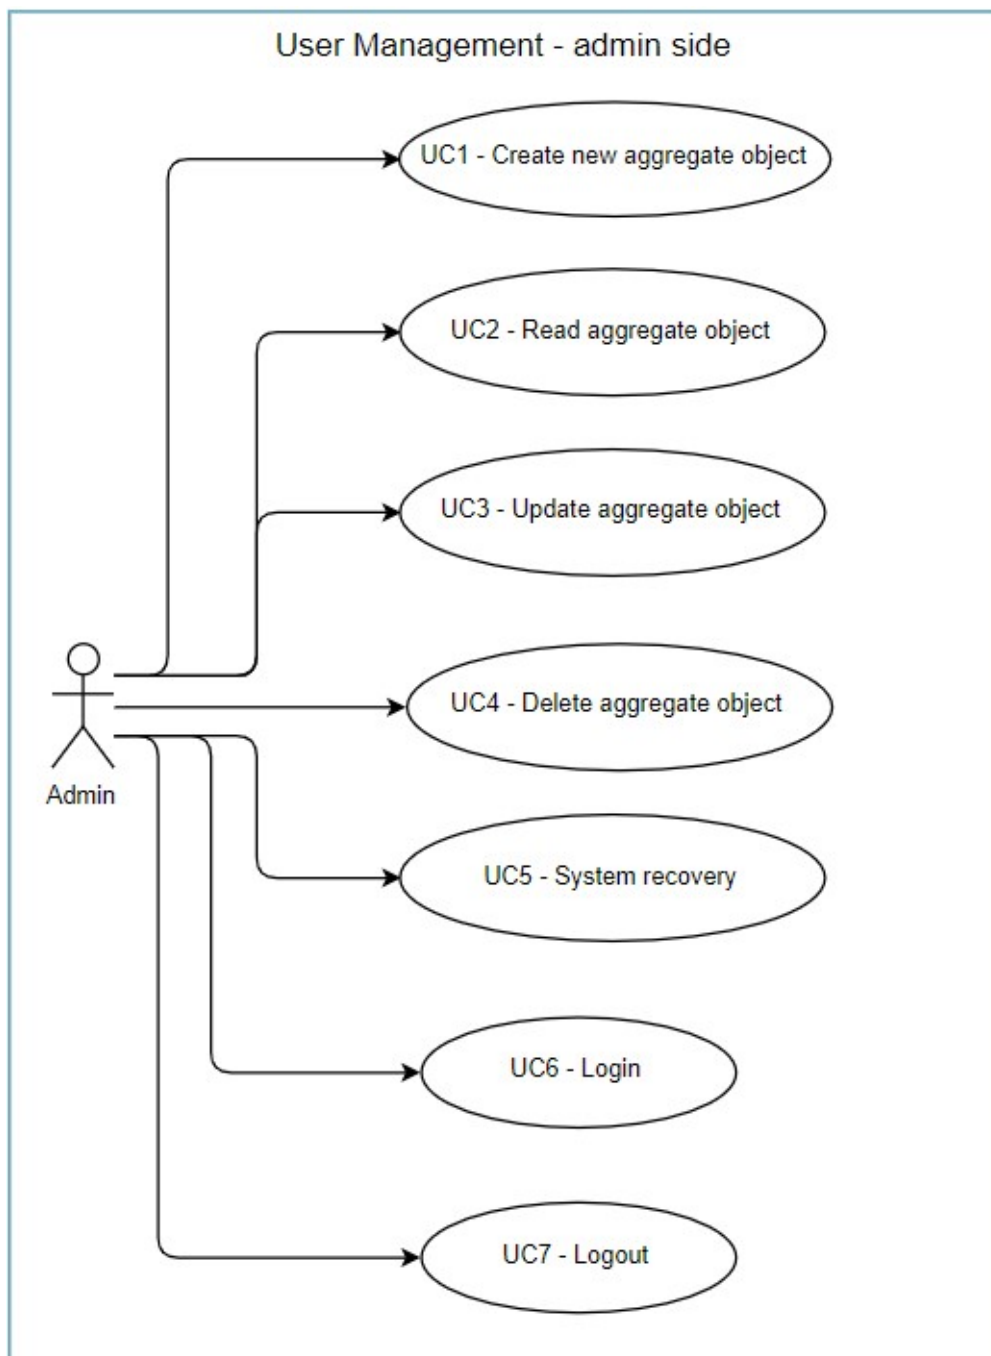


Fig. 12: Admin use cases

- UC1: the admin creates a new object based on the type of aggregate;
- UC2: the admin read the information about an object based on the type of aggregate;
- UC3: the admin updates an object based on the type of aggregate;
- UC4: the admin deletes an object;
- UC5: the admin recovers the system state from the chosen timestamp re-running of the event occurred after that;
- UC6: the admin logs in into the application using the Auth0 portal;
- UC7: the admin logs out from the application.

3.3 User side

3.3.1 Authentication

The authentication feature si provided by a third-party provider, Auth0. In the user side only users can log in; he can sign in with email and password or using Google's authentication or just logs in using his credentials.

3.3.2 Use cases




Fig. 13: User use cases

- UC1: the user fills the form to sign in into the application the first time logs in;
- UC2: the user read his profile;
- UC3: the user updates his profile;
- UC4: the user signs up to Auth0 authentication portal using email and password or Google's account;

- UC5: the user logs in to Auth0 authentication portal using email and password or Google's account;
- UC6: the user logs out from the application.

4 Serverless Framework

4.1 Description

The Serverless Framework is a free and open-source web framework written using Node.js. Serverless is the first framework that was originally developed for building applications exclusively on AWS Lambda, a serverless computing platform provided by Amazon as a part of the Amazon Web Services.

4.2 Serverless.yml

One advantage of using this tool is the capability to deploy every time an entire serverless system based on cloud providers, in this case AWS. You can define a *serverless.yml* configuration file which contains all the information about your service. You don't need to manually create the resources you need in a project, like database tables, queues, API endpoints and Lambda functions. You just have to write the code and then deploy your app.

4.2.1 Serverless configuration for user management

This snippet of serverless configuration shows the name of your service, a list of plugins and the custom field. The *dotenv* plugin is helpful if you have variables stored in a *.env* file that you want loaded into your serverless yaml config. This will allow you to reference them inside your config and it will load them into your lambdas. Into custom field you have to declare the name of your variables.

The *split-stack* plugin is useful when you reach the limit of 200 resources to deploy because it automatically splits your resources using nested stacks.

```
service: serverless-user-management
plugins:
  - serverless-dotenv-plugin
  - serverless-plugin-split-stacks #to avoid the limit of 200 resources
custom:
  splitStacks:
    perFunction: false
    perType: true
  dotenv:
    include:
      - AUTHO_ADMIN_CLIENT_ID
      - AUTHO_ADMIN_CLIENT_PUBLIC_KEY
      - AUTHO_ADMIN_DOMAIN
      - AUTHO_USER_CLIENT_ID
      - AUTHO_USER_CLIENT_PUBLIC_KEY
      - AUTHO_USER_DOMAIN
```

Fig. 14: Service, plugins and custom - serverless.yml

About your cloud provider: you can also define your IAM role authorization policies.

```
provider:
  name: aws
  runtime: nodejs10.x
  region: eu-central-1
  stage: dev
  iamRoleStatements:
    - Effect: "Allow"
      Resource: "*"
      Action:
        - "dynamodb:*"
        - "sqs:*"
        - "lambda:*"
        - "cloudwatch:*"
        - "apigateway:*
```

Fig. 15: Provider - serverless.yml

You can define custom error responses from API Gateway, helpful to avoid CORS error.

```
Failure500GatewayResponse:
  Type: 'AWS::ApiGateway::GatewayResponse'
  Properties:
    ResponseParameters:
      gatewayresponse.header.Access-Control-Allow-Origin: "'*'"
      gatewayresponse.header.Access-Control-Allow-Headers: "'*'"
    ResponseType: DEFAULT_5XX
    RestApiId:
      Ref: 'ApiGatewayRestApi'
```

Fig. 16: Custom gateway response - serverless.yml

DynamoDB tables parameters: you have to define the name of the table, the name and the type of the key and the provisioned throughput.

```
UsersDynamoDBTable:
  Type: 'AWS::DynamoDB::Table'
  Properties:
    AttributeDefinitions:
      - AttributeName: userId
    AttributeType: S
    KeySchema:
      - AttributeName: userId
    KeyType: HASH
    ProvisionedThroughput:
      ReadCapacityUnits: 5
      WriteCapacityUnits: 5
    TableName: user
```

Fig. 17: DynamoDB table - serverless.yml

Lambda function and its handler: you can define which events trigger that lambda. In this case the event triggers the function when a new object arrives into the specified queue.

```
commandCreateUser:
  handler: handler.commandCreateUser
  timeout: 10
  events:
    - sqs:
        arn: arn:aws:sqs:eu-central-1:582373673306:createUserQueue
```

Fig. 18: Lambda triggered by SQS event - serverless.yml

You can define an API gateway endpoint that triggers the corresponding function: you can also use a custom request template. If you want to protect the public endpoint you have to specify the authorizer function.

```
readUser:
  handler: handler.readUser
  events:
    - http:
        path: /readUser
        method: get
        authorizer:
          arn:
            arn:aws:lambda:eu-central-1:582373673306:function:serverless-user-man
          type: request
        resultTtlInSeconds: 0
  cors: true
```

```
integration: lambda
request:
  template:
    application/json: '{ "userId": "$input.params("userId")" }'
```

Fig. 19: Lambda triggered by GET endpoint - serverless.yml

4.3 Project's root

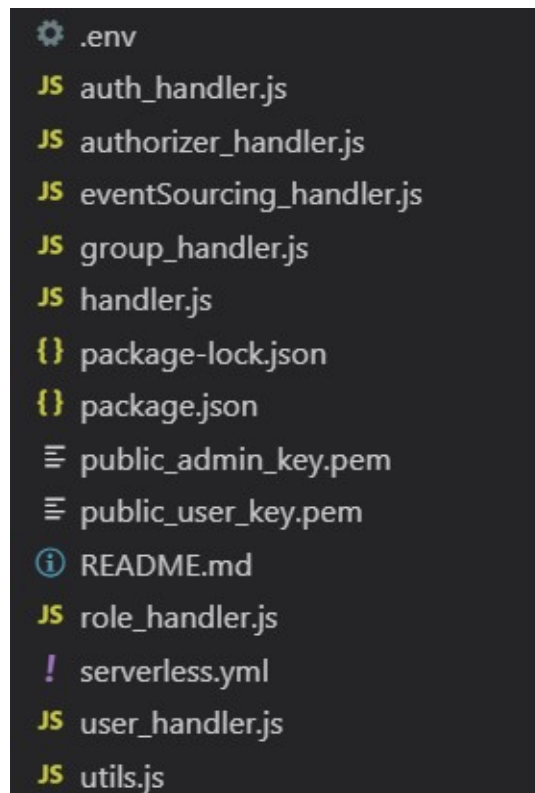


Fig. 20: Project's root

- The main handler imports other handler's files: one per aggregate, one for authorizers and one to handle event sourcing functions. In this way the responsibilities are restricted to every type of aggregate;
- The *utils.js* file contains a bunch of utilities functions;
- The *.env* file contains the environment variables loaded by the *dotenv* plugin;
- The *.pem* files contain the certificate provided by *Auth0* to authenticate the client;
- The *serverless.yml* contains the project's info and resources definition.

5 CQRS

Command Query Responsibility Segregation (CQRS) is an architectural pattern which separates the responsibility for modifying data (Command) from reading them (Query). The use of two different models for writing and reading operations allows to design and optimize each model for its responsibilities. The use of distinct models also allows the selection of the most appropriate technologies. As soon as the reading and writing models are separated, the infrastructure could easily scale to best fit the needs. It often happens that the number of writings in a system is much lower than the readings. Obviously the two models must be synchronized to ensure that the read information are consistent with the written ones.

The justification for CQRS is that in complex domains, a single model to handle both reads and writes gets too complicated, and we can simplify by separating the models.

The change that CQRS introduces is to split that conceptual model into separate models for update and display, which it refers to as Command and Query.

CQRS fits well with event-based programming models. It's common to see CQRS system split into separate services communicating with Event Collaboration. This allows these services to easily take advantage of Event Sourcing.

5.1 Architecture overview




Fig. 21: Architecture overview

As you can see from the architecture overview, the application is separated into two models: write model(right side) and read model(left side).

The choice to apply the CQRS pattern due to separate the responsibilities of writing and reading, but also because read side events aren't stored into the event store table because they don't change the system's state but just retrieve information from it.

5.2 Write model




Fig. 22: Write model sequence diagram

5.3 Read model




Fig. 23: Read model sequence diagram

6 Event Sourcing

"Event Sourcing ensures that all changes to application are stored as a sequence of events."

You can query the events, use the event log to reconstruct past states and adjust the state to cope with retroactive changes.

The core idea of event sourcing is that whenever we make a change to the state of a system, we record that state change as an event, and we can confidently rebuild the system state by reprocessing the events at any time in the future.

Before storing the event you must validate it because you don't have to store events that generate errors or exceptions.

When working with an event log can be useful to build snapshots of the working copy so that you don't have to process all the events when you rebuild the system or every time you need to query the database. For this reason when a new event is stored triggers another function to update the views.

Event sourcing can be used to:

- **Complete Rebuild:** you can discard the application state completely and rebuild it by re-running the events from the event store on an empty application.
- **Temporal Query:** you can determine the application state at any point in time. This can be used considering multiple time-lines (like branching in a VCS).
- **Event Replay:** if you find a past event was incorrect, you can replaying from then with the new event. The same technique can handle events received in the wrong sequence with systems that communicate with asynchronous messaging.

6.1 Event store

Event store is a database's table that contains all the events occurred from the begin. With event sourcing the event store becomes the principal source of truth and the system state is completely derived from it.

Is essential that the table can't be modified and the events that are replayed don't have to be stored again. This because the event store can reach huge dimensions and we must avoid to overload it.

6.1.1 Event object

Each event stored in the table is a JSON object and respect this format:

- **EventId:** is a UUID;
- **Payload:** JSON object which contains all the information about the event;
- **Aggregate:** type of the aggregate that the event refers to;
- **ExecutionQueue:** name of the execution queue used to execute the event;
- **Timestamp:** number to know the execution order of the events.

```
aggregate String : user
eventId String : f7731a40-f634-354a-ac99-351074d83b6d
executionQueue String : executeCreateUserQueue
► payload Map {7}
timestamp Number : 1563543008991
```

Fig. 24: Event object item

7 Extension points

7.1 New aggregates

To create a new type of aggregate you have to:

- Define a new DynamoDB table in the *serverless.yml*;
- Create an *operationQueue* and an *executionQueue* with SQS;
- Define the Lambda functions you need in the *serverless.yml*;
- Define the trigger events in the *serverless.yml*;
- Create a handler file named *"aggregate_handler.js"* in the project's root which contains the Lambda's handlers that refer the same aggregate type;
- Add the new handlers file in the *module.exports* of the main handler;
- Deploy the serverless application.

7.2 New write operation

To add a new write operation you have to:

- Create an *operationAggregateQueue* and an *executeOperationAggregateQueue*;
- Define a *pushOperationAggregateToSQS* with a POST API endpoint and the corresponding authorizer function to protect it in the *serverless.yml*;
- Define a *commandOperationAggregate* and an *operationAggregate* functions in the *serverless.yml*;
- Define trigger events in the *serverless.yml*;
- In the *"aggregate_handler.js"* write:
 - a *pushOperationAggregateToSQS* function to push the event in the corresponding queue;
 - a *commandOperationAggregate* function to check if the event is valid and store it into the *eventStore*;
 - an *operationAggregate* function to execute the event and update the view.
- Deploy the serverless application.

7.3 New read operation

To add a new read operation you have to:

- Define the function's name and the corresponding handler's name in the *serverless.yml*;
- Define the corresponding GET API endpoint in the *serverless.yml*;
- Define the authorizer function to protect the endpoint in the *serverless.yml*;

- In the *"aggregate_handler.js"* write the *readOperation* function to query the database;
- Deploy the serverless application.

8 Setup

8.1 Installing Node.js

Serverless Framework runs on Node v6 or higher so you have to install Node.js on your machine. You can download it from this link: <https://nodejs.org/it/download/>

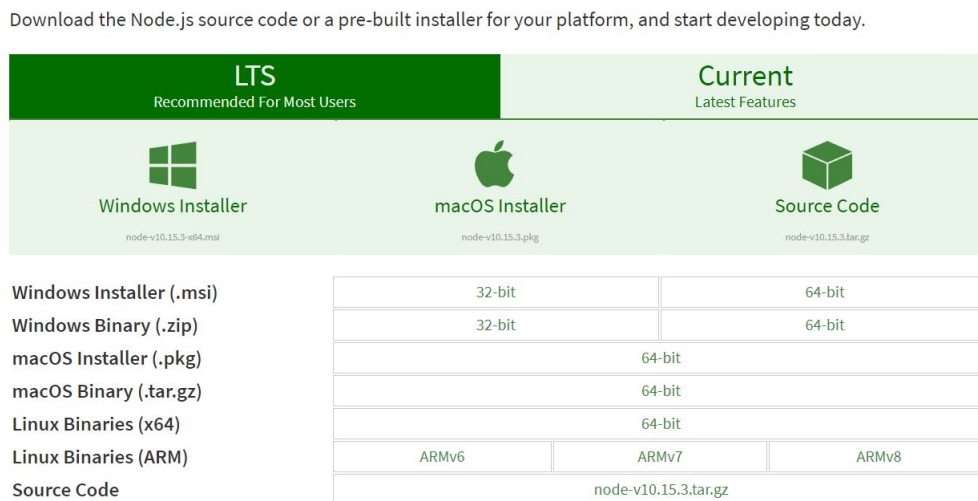


Fig. 25: Node.js download page

You can verify the correct installation running the commands: `node -v` and `npm -v`

8.2 Installing Serverless Framework

You can install the Serverless Framework running the command: `npm install -g serverless`

You can verify the correct installation of it running the command: `serverless -v` or `sls -v`

8.3 Setup AWS credentials

You have to configure your serverless CLI with yours AWS credentials, so you have to create an IAM policy related with the Serverless Framework:

1. Login to your AWS account and go to the Identity & Access Management(IAM) page;
2. Click on Users and then Add user;
3. Enter a name in the first field to remind you this User is related to the Serverless Framework;
4. Enable Programmatic access;

Set user details

You can add multiple users at once with the same access type and permissions. [Learn more](#)

User name*

[+ Add another user](#)

Select AWS access type

Select how these users will access AWS. Access keys and autogenerated passwords are provided in the last step. [Learn more](#)

Access type* ☒ **Programmatic access**
Enables an **access key ID** and **secret access key** for the AWS API, CLI, SDK, and other development tools.

☐ **AWS Management Console access**
Enables a **password** that allows users to sign-in to the AWS Management Console.

Fig. 26: Setup user details - AWS credentials

- Click on Attach existing policies directly, search for and select *AdministratorAccess*;

▼ Set permissions

Add user to group

Copy permissions from existing user

Attach existing policies directly

[Create policy](#) [Refresh](#)

Filter policies ▼ Showing 10 results

	Policy name ▼	Type	Used as	Description
<input checked="" type="checkbox"/>	AdministratorAccess	Job function	Permissions policy (3)	Provides full access to AWS services and...

Fig. 27: Setup policy - AWS credentials

- Create the user;
- Get the API key and secret;
- Run the command: `serverless config credentials --provider aws --key YOURKEY --secret YOURSECRET`

8.4 Deploying the serverless service

To deploy your serverless application you have to:

- Create a new directory;
- Create a new project running the command: `serverless` or `sls`
- Install dependencies running the command: `npm install`
- Install the plugins you need defining them into the `serverless.yml` and then running the command: `npm install -PLUGINNAME`
- Deploy your service running the command: `serverless deploy` or `sls deploy`