



Event-Driven cloud Architecture, CQRS and Event Sourcing for user management

2019/08/09

About

Company	Molo17 S.r.l.
Author	Lorenzo Busin
State	Approved
Use	Internal
Email	lorenzo.busin@gmail.com

Description

This document contains doc about cloud architectures that use an Event-Driven approach and implement CQRS and Event Sourcing for user management.

Table of contents

1	Introduction	3
1.1	Project purpose	3
1.2	References	3
2	Amazon Web Services	4
2.1	Lambda	4
2.1.1	Write mode Lambda functions	4
2.1.1.1	PushOperationAggregateToSQS	4
2.1.1.2	CommandOperationAggregate	5
2.1.1.3	Mediator	7
2.1.1.4	OperationAggregate	8
2.1.1.5	Recovery	9
2.1.2	Read mode Lambda functions	9
2.1.2.1	ReadOperationAggregate	9
2.1.3	CloudWatch Logs	10
2.2	DynamoDB	10
2.3	API Gateway	11
2.4	Simple Queue Service	11
3	Serverless Framework	12
3.1	Description	12
3.2	serverless.yml	12
4	CQRS	13
4.1	Architecture overview	13
4.2	Write model	13
4.3	Read model	13
5	Aggregates	14
5.1	User	14
5.2	Role	14
5.3	Authorization	14
5.4	Group	14
6	Event Sourcing	15
6.1	Event store	15
7	Extension points	16
7.1	New aggregates	16
7.2	New write functions	16
7.3	New read functions	16
8	Setup	17

List of figures

1	Fetch POST API	4
2	DynamoDB event object	7
3	Parsed DynamoDB event object	8
4	Fetch GET API	9

1 Introduction

1.1 Project purpose

This project explains how to build Event-Driven architectures, CQRS and Event Sourcing for user management, showing how should be implemented, extended, strengths and weaknesses. The application for user management implements an authentication function and the CRUD operations for each aggregates, which are: users, roles, authorizations and groups.

1.2 References

- What do you mean by "Event-Driven"? - Martin Fowler:
martinfowler.com/articles/201701-event-driven.html;
- CQRS - Martin Fowler:
martinfowler.com/bliki/CQRS.html;
- Event Sourcing - Martin Fowler:
martinfowler.com/eaaDev/EventSourcing.html;
- AWS Lambda documentation:
<https://docs.aws.amazon.com/lambda/index.html>;
- AWS DynamoDB documentation:
<https://docs.aws.amazon.com/dynamodb/index.html>;
- AWS API gateway documentation:
<https://docs.aws.amazon.com/apigateway/index.html>;
- AWS SQS documentation:
<https://docs.aws.amazon.com/sqs/index.html>;

2 Amazon Web Services

Amazon Web Services (AWS) is a subsidiary of Amazon that provides on-demand cloud computing platforms to individuals, companies, and governments, on a pay-as-you-go basis. These cloud computing web services provide a set of primitive abstract technical infrastructure and distributed computing building blocks and tools. AWS's version of virtual computers emulate most of the attributes of a real computer including, hardware central processing units and graphics processing units, local/RAM memory, hard-disk/SSD storage; a choice of operating systems; networking; and pre-loaded application software such as web servers and databases.

2.1 Lambda

AWS Lambda is an event-driven, serverless computing platform provided by Amazon. It is a computing service that runs code in response to events and automatically manages the computing resources required by that code. The purpose of Lambda is to simplify building smaller, on-demand applications that are responsive to events and new information. AWS targets starting a Lambda instance within milliseconds of an event. Node.js, Python, Java, Go, Ruby and C# through .NET Core are all officially supported.

In this project, lambda functions are the computing core for the execution of all commands and are written in Node.js.

2.1.1 Write mode Lambda functions

2.1.1.1 PushOperationAggregateToSQS

Those functions starts the flow and triggered when fetching the corresponding API gateway URL with a POST request like this:

```
fetch("https://domain/resourceAPIpath", {
  method: "post",
  headers: {
    'Accept': 'application/json',
    'Content-Type': 'application/json'
  },

  body: JSON.stringify({
    "attribute1": "value1",
    "attribute2": "value2"
  })
});
```

Figura 1: Fetch POST API

Once triggered, these functions retrieve the event and put it in the *MessageBody* parameter. The event object is not changed, instead in the user functions which encrypt the password before sending the message to corresponding SQS queue.

Example:

```

1 module.exports.pushUpdateRoleToSQS = (event, context, callback) => {
2   const AWS = require('aws-sdk');
3   const SQS = new AWS.SQS();
4   const stringedEvent = JSON.stringify(event);
5
6   const params = {
7     MessageBody: stringedEvent,
8     QueueUrl: "https://yourDomain/updateRoleQueue"
9   };
10
11  SQS.sendMessage(params, function(err, data){
12    if(err){
13      console.log(err);
14      callback(null, err);
15    }
16    else
17      callback(null, "Role event pushed to SQS");
18  });
19 };

```

2.1.1.2 CommandOperationAggregate

These functions validate the values of the attributes before storing the event in the *eventStore* and triggered when a new event arrives to the corresponding queue.

If you have to check for duplicated attributes in the database, the function must be marked *async* because it has to wait for the result of the check operation.

Example:

```

1 module.exports.commandCreateRole = async (event, context, callback) => {
2   const utils = require('./utils.js');
3
4   const stringedEvent = event.Records[0].body.toString('utf-8');
5   const eventParsed = JSON.parse(stringedEvent);
6   const stringedBody = JSON.stringify(eventParsed);
7   const bodyParsed = JSON.parse(stringedBody);
8   const check = bodyParsed.body;
9
10  const checkNameParams = {
11    TableName: 'role',
12    ExpressionAttributeNames: {
13      "#rolename": "name"
14    },
15    ProjectionExpression: "#rolename",
16    FilterExpression: "#rolename = :checkname",
17    ExpressionAttributeValues: {

```

```
18         ":checkname": check.name
19     }
20
21     const nameAlreadyExists = await utils.asyncCheckScanDB(checkNameParams);
22
23     if ((check.roleId == "" || check.name == "" || check.desc == "") || ↵
24         nameAlreadyExists)
25         callback(null, "Name already exists or empty attributes");
26     else{
27         utils.storeEvent("role", "executeCreateRoleQueue", bodyParsed.body);
28         callback(null, "Role event stored");
29     }
30 };
```

2.1.1.3 Mediator

This function catch the *DynamoDB* event and triggered when a change occurs in a certain table.

You have to be careful that the *DynamoDB* events are mapped using the char type attribute value. This is an example:

```
{
  "eventId": {
    "S": "bf6dffb9-72d8-ae5f-21fa-56dd6a26d572"
  },
  "payload": {
    "M": {
      "auth": {
        "S": "{\n\t\"Authorizations\": [\n\t\t\"FullAccess\" ] }"
      },
      "roleId": {
        "S": "63c471c1-e5c7-09d0-ea8a-f20b27a0575c"
      },
      "name": {
        "S": "Admin"
      },
      "desc": {
        "S": "Full access to all resources"
      }
    }
  },
  "aggregate": {
    "S": "role"
  },
  "executionQueue": {
    "S": "executeCreateRoleQueue"
  },
  "timestamp": {
    "N": "1563358792450"
  }
}
```

Figura 2: DynamoDB event object

You can use the *"AWS.DynamoDB.Converter"* module to parse a DynamoDB event object.


```

{
  "eventId": "bf6dffb9-72d8-ae5f-21fa-56dd6a26d572",
  "payload": {
    "auth": "{\n\t\"Authorizations\": [\n\t\t\"FullAccess\" ] }",
    "roleId": "63c471c1-e5c7-09d0-ea8a-f20b27a0575c",
    "name": "Admin",
    "desc": "Full access to all resources"
  },
  "aggregate": "role",
  "executionQueue": "executeCreateRoleQueue",
  "timestamp": 1563358792450
}

```

Figura 3: Parsed DynamoDB event object

After that, the mediator retrieves the *executionQueue* parameter from the event object, the payload event passed with the *MessageBody* and then sends the execution message to corresponding SQS queue.

2.1.1.4 OperationAggregate

These functions execute a single operation using the event payload and triggered when new event arrives to the corresponding execution queue.

Example:

```

1 module.exports.createRole = async (event, context, callback) => {
2   const AWS = require('aws-sdk');
3   const dynamoDb = new AWS.DynamoDB.DocumentClient();
4
5   const stringedBody = event.Records[0].body.toString('utf-8');
6   const parsedBody = JSON.parse(stringedBody);
7
8   const params = {
9     TableName: 'role',
10    Item: parsedBody
11  };
12
13  await dynamoDb.put(params, (err, data) => {
14    if (err){
15      console.log(err);
16      callback(null, err);
17    }
18    else
19      callback(null, "Role created");
20  }).promise();
21 };

```

2.1.1.5 Recovery

This function allows you to restore the system state starting from a given timestamp by replaying all the events into the *eventStore* table which were stored after that time.

Is important to re-execute the events one by one and in the correct order(from the oldest one); for this purpose the recovery function implements a sorting algorithm that retrieves an array of events and sorts them by timestamp, and an *async* function which sends every event to the corresponding execution queue.

2.1.2 Read mode Lambda functions

2.1.2.1 ReadOperationAggregate

These functions works in the read side of the architecture; they query the database to retrieve the informations needed without passing throw a SQS queue or a mediator. This kind of events aren't stored into the *eventStore* because they don't change the current system state; this Lambda functions fetch the URL of a GET endpoint using a GET request and listen for a response result.

```
fetch("https://domain/resourceAPIpath" [+ queryStringParam] {  
  method: "GET",  
  headers: {  
    'Content-Type': 'application/json'  
  }  
});
```

Figura 4: Fetch GET API

Example:

```
1 module.exports.getAllRoles = (event, context, callback) => {  
2   const AWS = require('aws-sdk');  
3   const dynamoDb = new AWS.DynamoDB.DocumentClient();  
4  
5   const params = {  
6     TableName: 'role',  
7     ExpressionAttributeNames: {  
8       "#rolename": "name"  
9     },  
10    ProjectionExpression: "#rolename"  
11  };  
12  
13  dynamoDb.scan(params, (err, data) => {  
14    const stringedData = JSON.stringify(data);  
15    if (err) {  
16      console.log(err);  
17      callback(null, err);  
18    }  
19  });
```

```
19     else{
20         if (data.Count == 0){
21             console.log("Role not found");
22             callback(null, "Role not found");
23         }
24         else{
25             const response = {
26                 statusCode: 200,
27                 headers: {
28                     'Content-Type': 'application/json',
29                     'Access-Control-Allow-Origin': '*',
30                     'Access-Control-Allow-Credentials': true
31                 },
32                 body: stringedData
33             };
34             callback(null, response);
35         }
36     }
37 });
38 };
```

2.1.3 CloudWatch Logs

Amazon CloudWatch is a monitoring and management service built for developers, system operators, site reliability engineers and IT managers. CloudWatch provides you with data and actionable insights to monitor your applications, understand and respond to system-wide performance changes, optimize resource utilization, and get a unified view of operational health. CloudWatch collects monitoring and operational data in the form of logs, metrics, and events.

This tool can be very helpful to debug your Lambda functions and to understand what happens to your system.

2.2 DynamoDB

Amazon DynamoDB is a fully managed proprietary NoSQL database service that supports key-value and document data structures and is offered by Amazon Web Services. In this project the DynamoDB instance is used for two purpose: it is used to store events and to keep updated the aggregates views.

The tables are the following:

- **eventStore**: to keep track of the occurred events;
- **user**: to store users info;
- **role**: to store roles info;
- **authorization**: to store authorizations info;
- **group**: to store groups info.

2.3 API Gateway

Amazon API Gateway is a fully managed service that makes it easy for developers to create, publish, maintain, monitor, and secure API at any scale. You can create REST and WebSocket API that act as a “front door” for applications to access data, business logic, or functionality from your backend services.

In this project, API endpoints starting the execution flow of a write mode function through a POST request, pushing a new event into the corresponding queue; in fact you have to provide an API endpoint for each function. On the other side, to query the database you have to reach the endpoint through a GET request to get a result response.

2.4 Simple Queue Service

Simple Queue Service(SQS) is a distributed message queuing service introduced by Amazon. It supports programmatic sending of messages via web service applications as a way to communicate over the Internet. SQS is intended to provide a highly scalable hosted message queue that resolves issues arising from the common producer-consumer problem or connectivity between producer and consumer.

In this project, for each function you have to use two queues:

- **OperationAggregateQueue:** this queues receive new events from an API endpoint and trigger the corresponding *commandOperationAggregate* function;
- **ExecuteOperationAggregateQueue:** this queues receive new events from the *mediator* function and trigger the corresponding *operationAggregate* function.

3 Serverless Framework

3.1 Description

3.2 `serverless.yml`

4 CQRS

4.1 Architecture overview

4.2 Write model

4.3 Read model

5 Aggregates

5.1 User

5.2 Role

5.3 Authorization

5.4 Group

6 Event Sourcing

6.1 Event store

7 Extension points

7.1 New aggregates

7.2 New write functions

7.3 New read functions

8 Setup